

Predicting Customer Turnover

And How We See It Coming


Tim McAleer

The Problem

- SyriaTel, a communications company, wants to identify customers who leave their services.
- Preferably, before they leave.
- Given a database of 3,333 customers detailing many characteristics include call details, payments, and whether or not that customer retained services, develop a model to identify customers who are likely to leave and note the significant identifiers.


The Obstacle

Class Imbalance

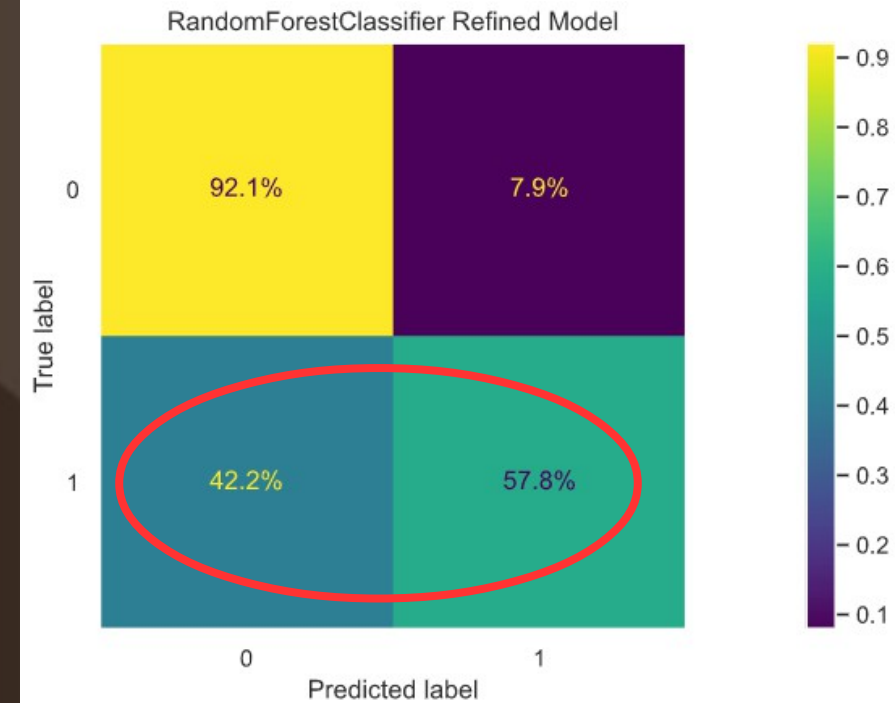

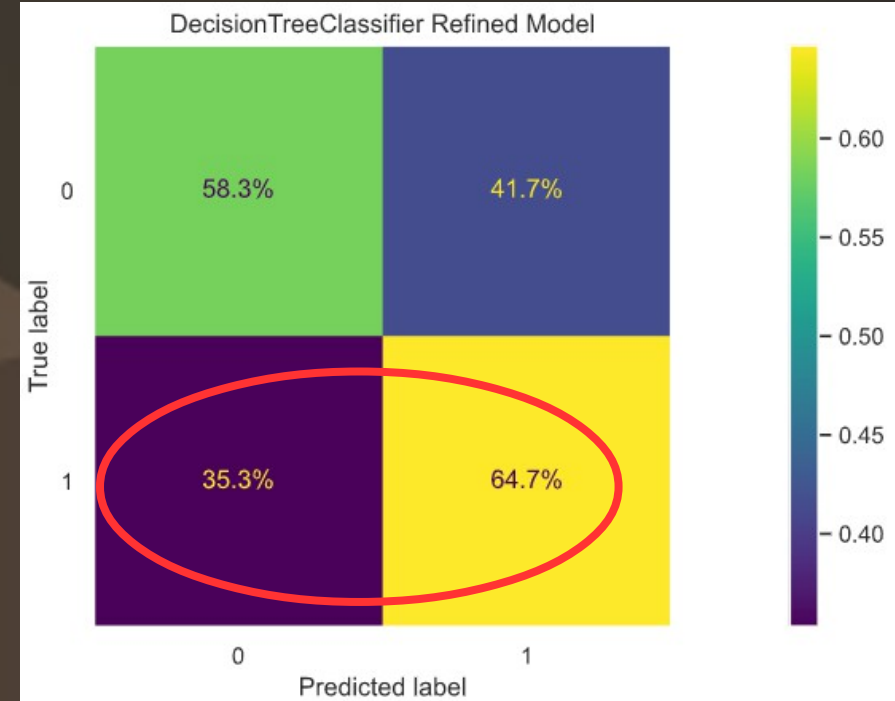
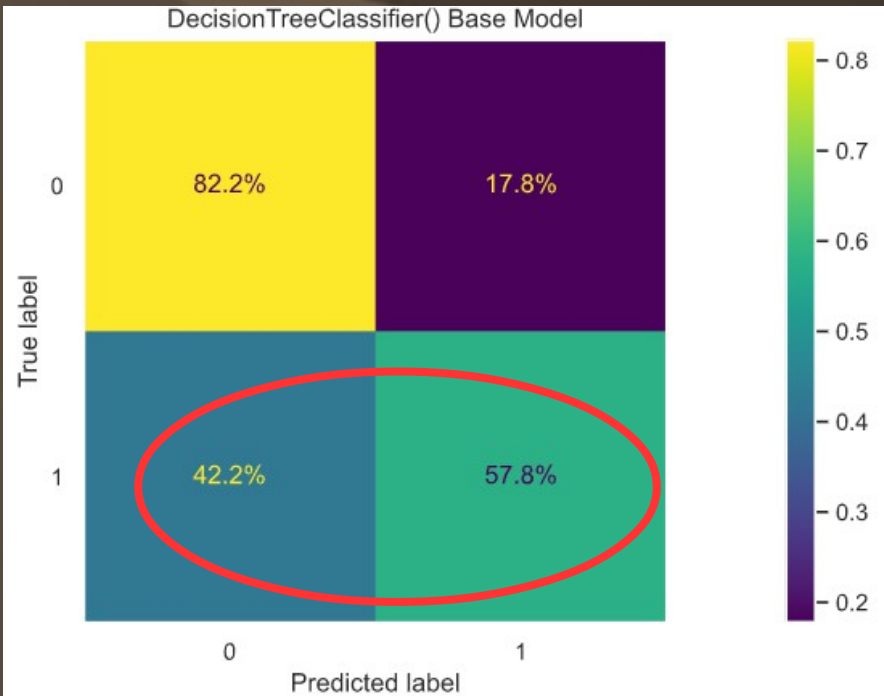


Dealing With Imbalance

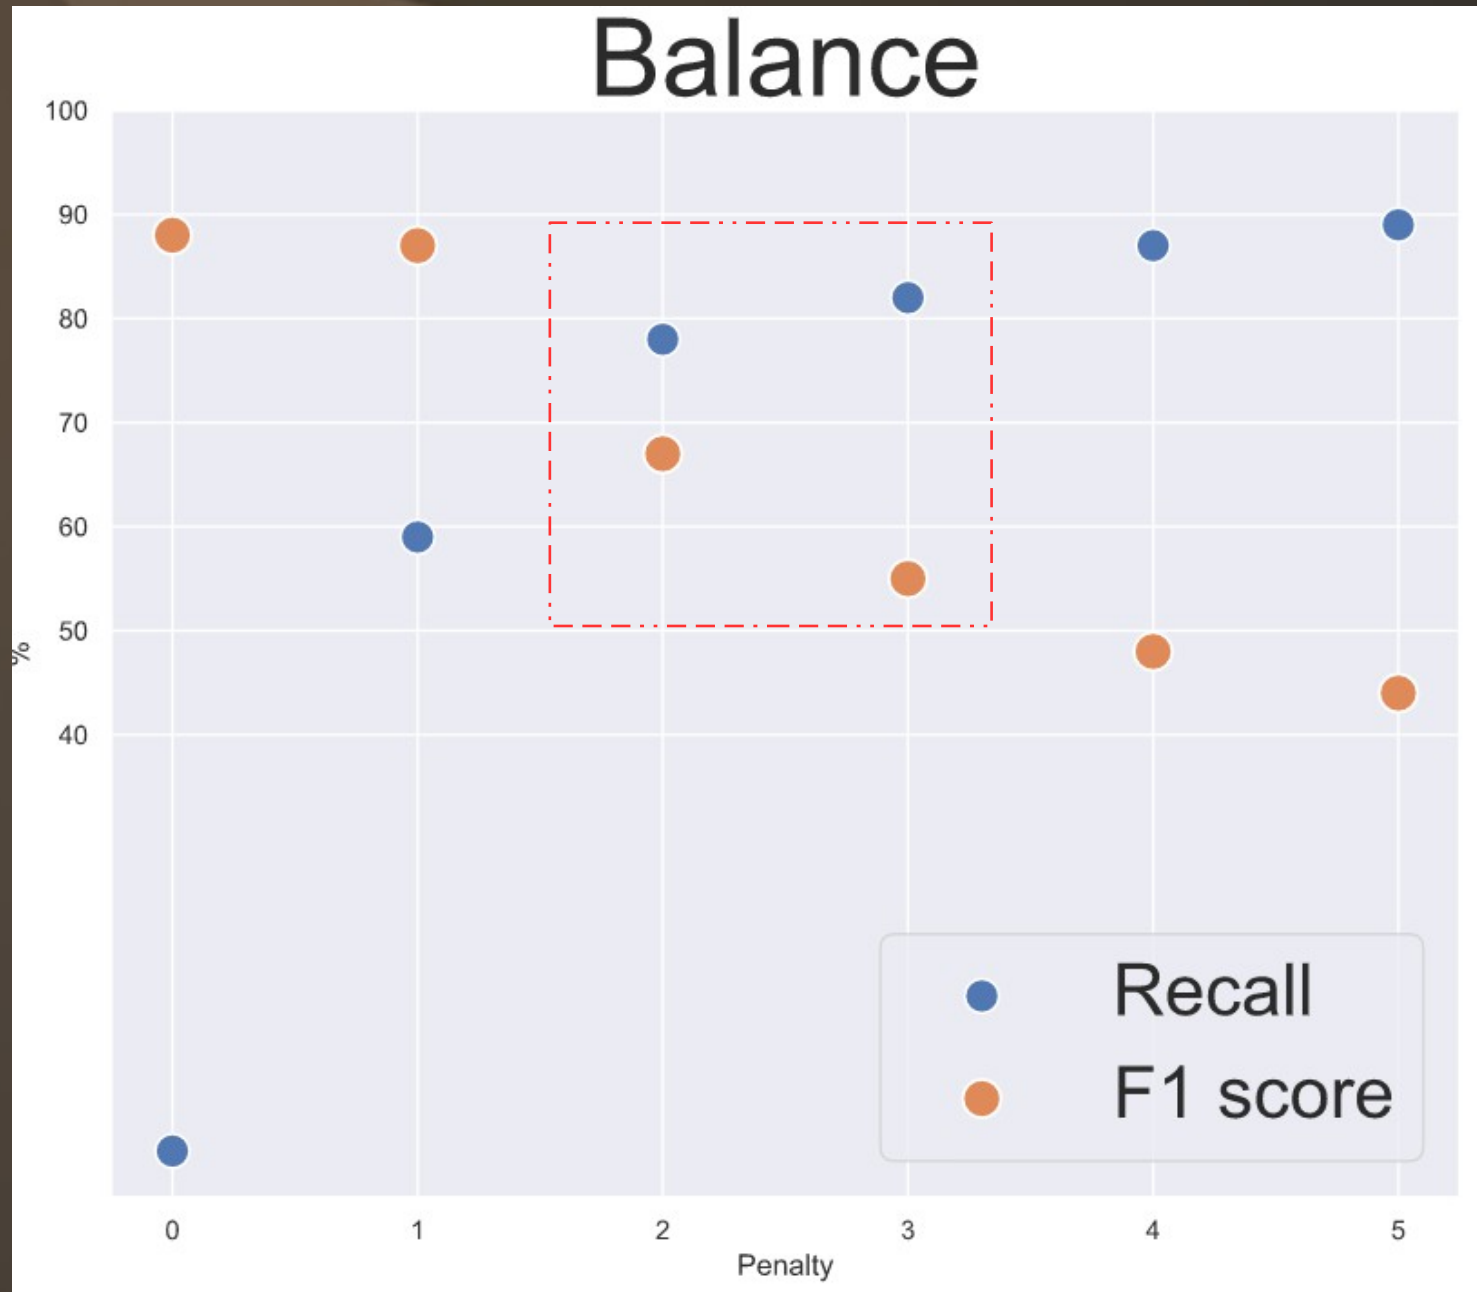
- Oversampling
- Undersampling
- VAE
- SMOTE




Choosing a Model



XGBoost – Recall vs F1



Customers Who Leave Have These Features



Thanks for listening!

- Contact me! I prefer compliments.
tcmcaleer@gmail.com
- Check out the work!
<https://github.com/tcmcaleer/SyriaTel-Customer-Retention>
- Thanks to Flatiron for their support!
- Thanks to all of you!

